

# ▼ 1 SyriaTel Customer Analysis

By: Brian Lee

---

## ▼ 1.1 Business problem

**SyriaTel** telecommunications company has hired us to analyze the causes to customer churn. **Churn** is whether a customer will stop doing business with the company. We are trying to predict the likelihood of a customer churn based on a user's communication usage, plans, and other related factors. SyriaTel can use these findings to improve the services to better keep customers and maintain greater profit for the company.

## ▼ 1.2 Data Understanding

This project uses the SyriaTel dataset, which can be found in 'telecoms.csv' in the 'data' folder.

```
In [1]: # Import necessary packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from statsmodels.formula.api import ols
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
import scipy.stats as stats
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.metrics import mean_absolute_error

import xgboost as xgb

from sklearn.metrics import roc_curve, auc

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB, BernoulliNB, MultinomialNB, ComplementNB
from sklearn.feature_selection import SelectKBest
import joblib

from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
from sklearn.metrics import confusion_matrix, recall_score, precision_recall_curve
from sklearn.metrics import precision_recall_fscore_support, f1_score, fbeta_score
from sklearn.metrics import classification_report, plot_roc_curve, plot_confusion_matrix
from sklearn.linear_model import LogisticRegression
from imblearn.over_sampling import SMOTE
from collections import Counter
from sklearn.metrics import classification_report

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier

from xgboost import XGBClassifier
```

executed in 1.97s, finished 00:56:16 2021-04-23

```
In [2]: df = pd.read_csv('data/telecoms.csv')
df.head()
```

executed in 47ms, finished 00:56:16 2021-04-23

Out[2]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total evening calls
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122

5 rows × 21 columns



▼ **1.3 Data Preparation**

Let's quickly examine the dataset and clean it up for proper analysis and modeling

```
In [3]: display(df.info())
display(df.describe())
```

executed in 77ms, finished 00:56:16 2021-04-23

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                  3333 non-null   object
5   voice mail plan                    3333 non-null   object
6   number vmail messages              3333 non-null   int64
7   total day minutes                   3333 non-null   float64
8   total day calls                     3333 non-null   int64
9   total day charge                    3333 non-null   float64
10  total eve minutes                   3333 non-null   float64
11  total eve calls                     3333 non-null   int64
12  total eve charge                    3333 non-null   float64
13  total night minutes                 3333 non-null   float64
14  total night calls                   3333 non-null   int64
15  total night charge                  3333 non-null   float64
16  total intl minutes                  3333 non-null   float64
17  total intl calls                    3333 non-null   int64
18  total intl charge                   3333 non-null   float64
19  customer service calls              3333 non-null   int64
20  churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

None

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes
<b>count</b>	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
<b>mean</b>	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348
<b>std</b>	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844
<b>min</b>	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000
<b>50%</b>	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000
<b>75%</b>	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000
<b>max</b>	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000

In [4]: *# Adjust column names with '\_'*

```
df.columns = df.columns.str.replace(' ', '_')
df.info()
```

executed in 12ms, finished 00:56:16 2021-04-23

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	state	3333 non-null	object
1	account_length	3333 non-null	int64
2	area_code	3333 non-null	int64
3	phone_number	3333 non-null	object
4	international_plan	3333 non-null	object
5	voice_mail_plan	3333 non-null	object
6	number_vmail_messages	3333 non-null	int64
7	total_day_minutes	3333 non-null	float64
8	total_day_calls	3333 non-null	int64
9	total_day_charge	3333 non-null	float64
10	total_eve_minutes	3333 non-null	float64
11	total_eve_calls	3333 non-null	int64
12	total_eve_charge	3333 non-null	float64
13	total_night_minutes	3333 non-null	float64
14	total_night_calls	3333 non-null	int64
15	total_night_charge	3333 non-null	float64
16	total_intl_minutes	3333 non-null	float64
17	total_intl_calls	3333 non-null	int64
18	total_intl_charge	3333 non-null	float64
19	customer_service_calls	3333 non-null	int64
20	churn	3333 non-null	bool

```
dtypes: bool(1), float64(8), int64(8), object(4)
```

```
memory usage: 524.2+ KB
```

```
In [5]: # Find missing values columns
df.isna().sum()
```

executed in 15ms, finished 00:56:16 2021-04-23

```
Out[5]: state                0
account_length             0
area_code                  0
phone_number               0
international_plan         0
voice_mail_plan            0
number_vmail_messages      0
total_day_minutes          0
total_day_calls            0
total_day_charge           0
total_eve_minutes          0
total_eve_calls            0
total_eve_charge           0
total_night_minutes        0
total_night_calls          0
total_night_charge         0
total_intl_minutes         0
total_intl_calls           0
total_intl_charge          0
customer_service_calls     0
churn                      0
dtype: int64
```

- No missing values found.

Phone numbers are just unique identifiers. Does not provide additional information. Could drop it

```
In [6]: # Drop phone_number
df.drop('phone_number', axis=1, inplace=True)
```

executed in 14ms, finished 00:56:16 2021-04-23

```
In [7]: # Values for area_code

df.area_code.value_counts()
```

executed in 15ms, finished 00:56:16 2021-04-23

```
Out[7]: 415    1655
510      840
408      838
Name: area_code, dtype: int64
```

Taking a look at the area\_code values show that there are only 3 area codes (San Francisco area), despite the data being declared for several different states. It will be better to drop the area codes in this case.

In [8]: *# Drop area\_code column*

```
df.drop('area_code', axis = 1, inplace=True)
```

executed in 14ms, finished 00:56:16 2021-04-23

Let's change the categorical columns to integers for easier analysis

In [9]: *# Change 'churn' from bool to int*

```
df['churn'] *= 1
```

*# Change plans to int*

```
df['international_plan'] = df['international_plan'].apply(lambda x: 1 if x=='yes' else 0)
df['voice_mail_plan'] = df['voice_mail_plan'].apply(lambda x: 1 if x=='yes' else 0)
```

```
df[['churn', 'international_plan', 'voice_mail_plan']].astype(int)
```

executed in 29ms, finished 00:56:16 2021-04-23

Out[9]:

	churn	international_plan	voice_mail_plan
0	0	0	1
1	0	0	1
2	0	0	0
3	0	1	0
4	0	1	0
...	...	...	...
3328	0	0	1
3329	0	0	0
3330	0	0	0
3331	0	1	0
3332	0	0	1

3333 rows × 3 columns

In [10]: `df.head()`

executed in 29ms, finished 00:56:16 2021-04-23

Out[10]:

	state	account_length	international_plan	voice_mail_plan	number_vmail_messages	total_day_mi
0	KS	128	0	1	25	
1	OH	107	0	1	26	
2	NJ	137	0	0	0	
3	OH	84	1	0	0	
4	OK	75	1	0	0	


## ▼ 1.4 Exploratory Data Analysis

Using the cleaned data, we will examine the distributions of the columns and descriptive statistics for the dataset



```
pd.plotting.scatter_matrix(df, figsize=[20,20]);  
plt.show()
```



executed in 16.2s, finished 00:56:33 2021-04-23

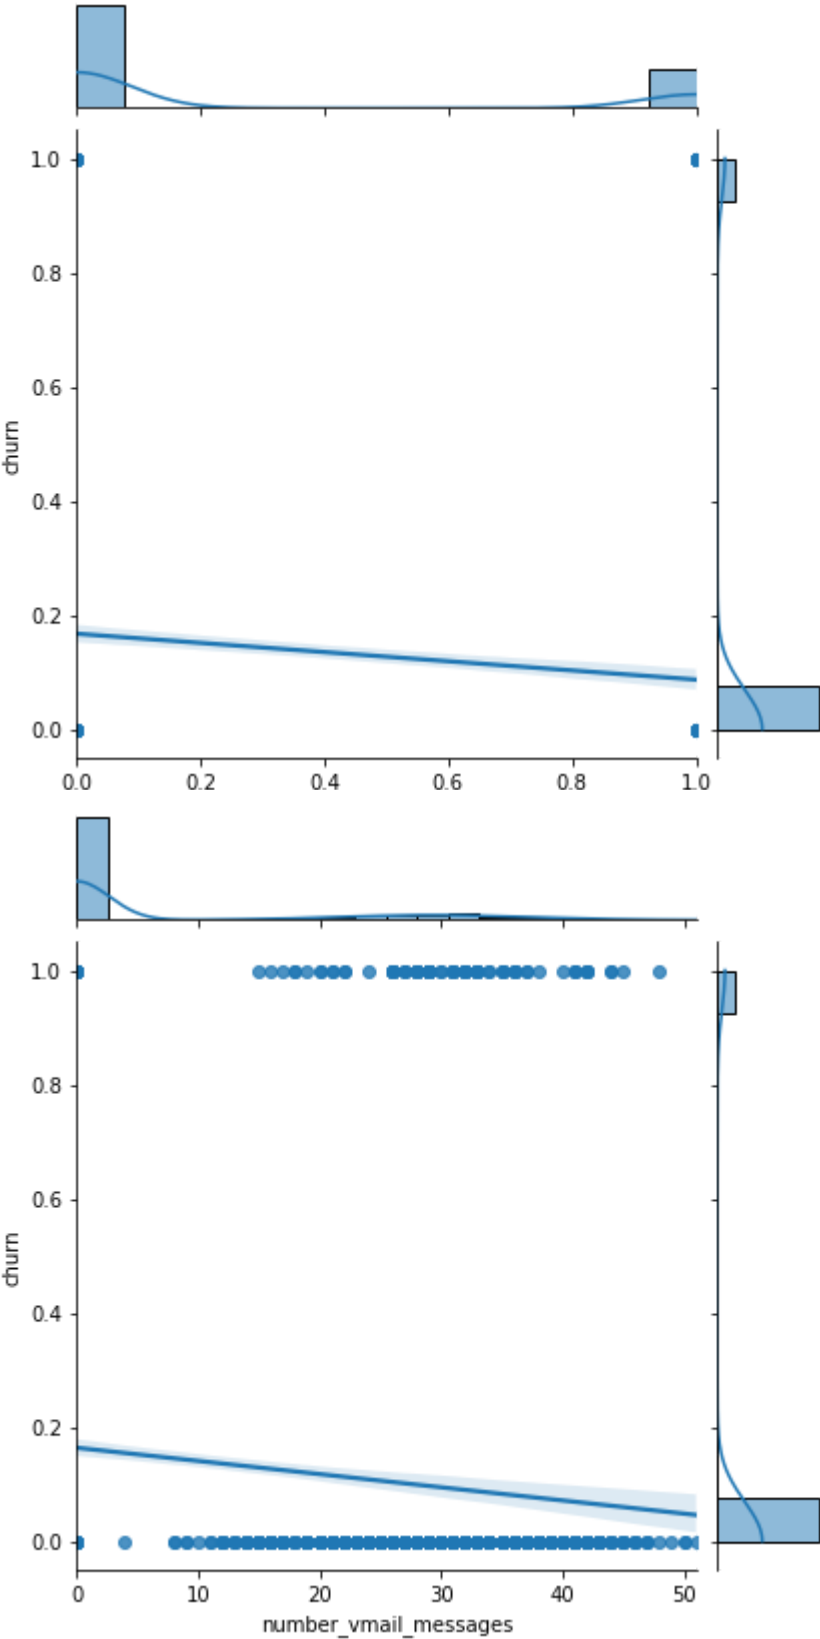


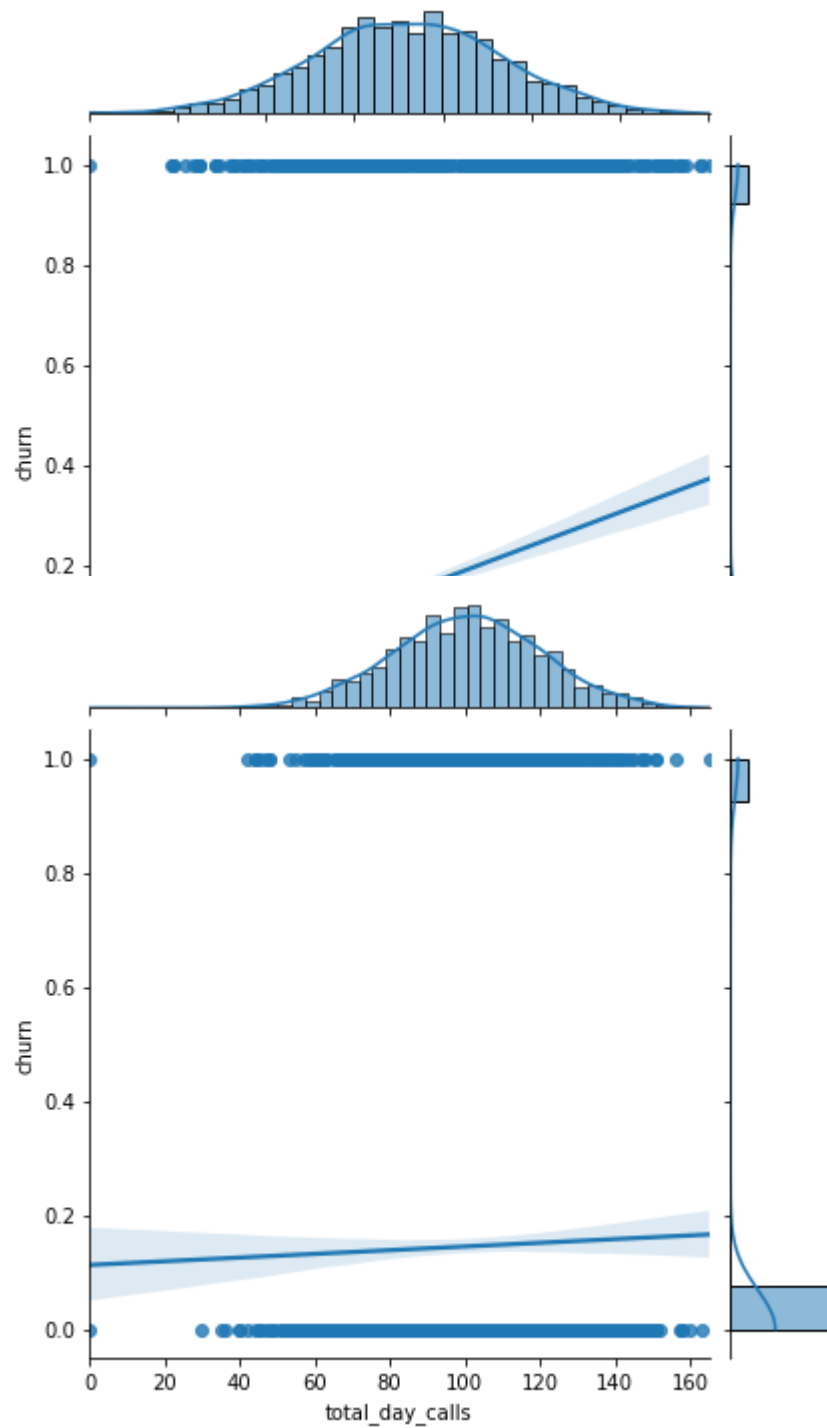




```
In [12]: for col in df.columns[1:]:  
         sns.jointplot(x=col, y='churn', data=df, kind='reg');
```


executed in 11.6s, finished 00:56:44 2021-04-23






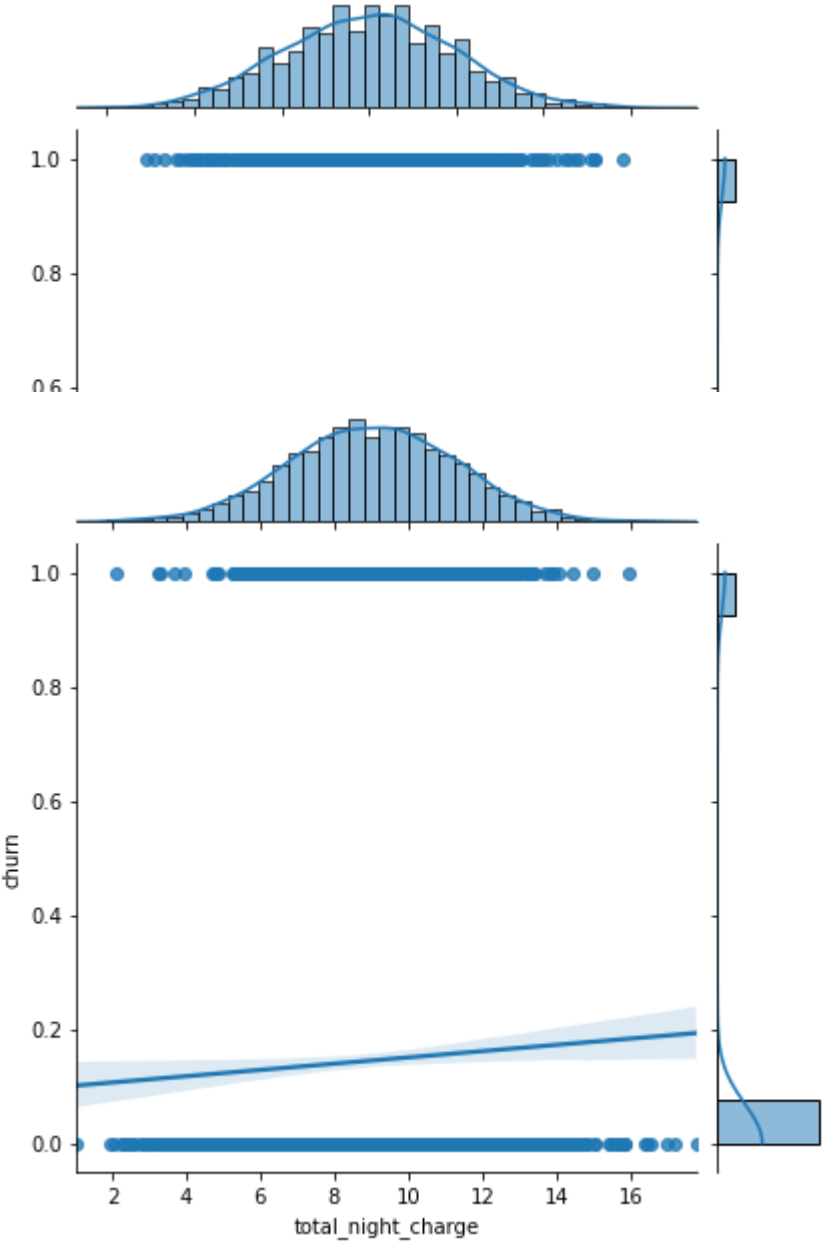
















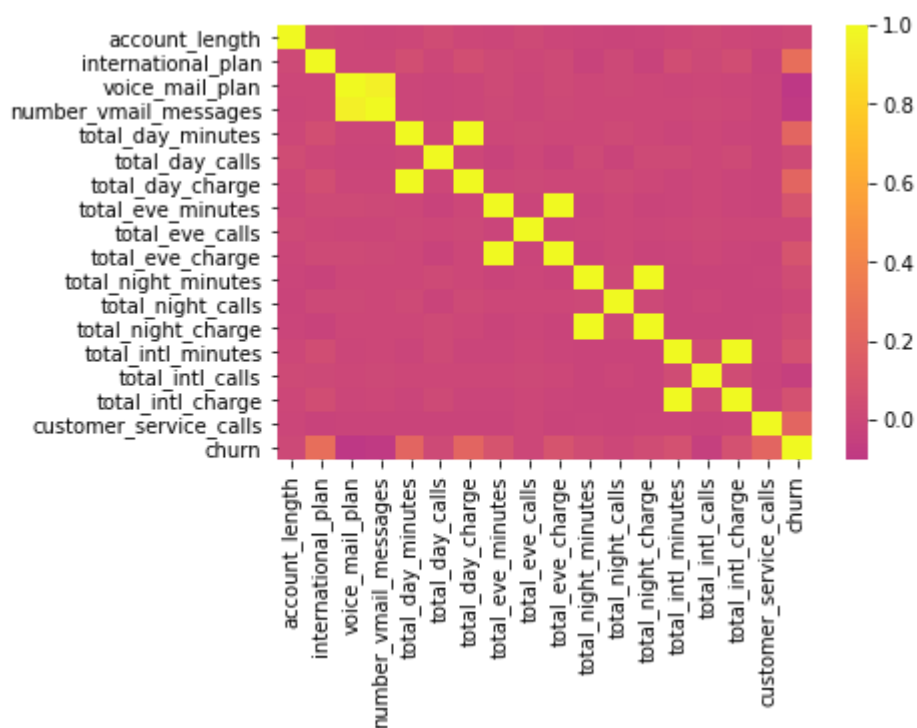






```
In [13]: # Heatmap for correlation values
import seaborn as sns
sns.heatmap(df.corr(), cmap='plasma', center=0);
```

executed in 486ms, finished 00:56:45 2021-04-23



There appears to be some heavy multicollinearity between some factors. Let's identify which ones are causing an issue

```
In [14]: # Displays if correlation coefficient values is greater than 0.75
df.corr()
abs(df.corr()) > 0.75

# Finds which column pairs have a CC values > 0.75
df_mc = df.corr().abs().stack().reset_index().sort_values(0, ascending=False)

df_mc['pairs'] = list(zip(df_mc.level_0, df_mc.level_1))

df_mc.set_index(['pairs'], inplace = True)

df_mc.drop(columns=['level_1', 'level_0'], inplace = True)

# cc for correlation coefficient
df_mc.columns = ['cc']

df_mc.drop_duplicates(inplace=True)

df_mc[(df_mc.cc>.75) & (df_mc.cc<1)]
```

executed in 30ms, finished 00:56:45 2021-04-23

Out[14]:

	cc
pairs	
(total_day_minutes, total_day_charge)	1.000000
(total_eve_minutes, total_eve_charge)	1.000000
(total_night_charge, total_night_minutes)	0.999999
(total_intl_charge, total_intl_minutes)	0.999993
(number_vmail_messages, voice_mail_plan)	0.956927

- Let's get rid of the factors related to minutes and keep the charge factors. Price is probably more important to our overall analysis
- We will also remove number\_vmail\_messages as the more important factor is that they have a voicemail plan.

```
In [15]: # Dropping total_day_minutes, total_eve_minutes, total_night_minutes, total_intl_
df.drop(['total_day_minutes', 'total_eve_minutes', 'total_night_minutes', 'total_intl_

```

executed in 15ms, finished 00:56:45 2021-04-23

In [16]: `df.head()`

executed in 14ms, finished 00:56:45 2021-04-23

Out[16]:

	state	account_length	international_plan	voice_mail_plan	total_day_calls	total_day_charge	total_eve_charge
0	KS	128	0	1	110	45.07	29.66
1	OH	107	0	1	123	27.47	18.29
2	NJ	137	0	0	114	41.38	27.31
3	OH	84	1	0	71	50.90	33.05
4	OK	75	1	0	113	28.34	18.29

## ▼ 1.5 Modeling

Now that we have explored the cleaned data, we can finally move on to create models to properly see the effects of each of the factors on telecom customer churning.

### ▼ 1.5.1 Logistic Regression (Baseline)

In [17]: `df_dummy = pd.get_dummies(df)`  
`df_dummy.head()`

executed in 31ms, finished 00:56:45 2021-04-23

Out[17]:

	account_length	international_plan	voice_mail_plan	total_day_calls	total_day_charge	total_eve_charge
0	128	0	1	110	45.07	29.66
1	107	0	1	123	27.47	18.29
2	137	0	0	114	41.38	27.31
3	84	1	0	71	50.90	33.05
4	75	1	0	113	28.34	18.29

5 rows × 64 columns

In [18]: `# Create X,y and train/test`  
`X = df_dummy.drop(columns=['churn'], axis=1)`  
`y = df_dummy['churn']`  
  
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state=42)`

executed in 14ms, finished 00:56:45 2021-04-23

```
In [19]: # initial regression
logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')
model_log = logreg.fit(X_train, y_train)
```

executed in 60ms, finished 00:56:45 2021-04-23

```
In [20]: # Prediction
y_hat_train = logreg.predict(X_train)
y_hat_test = logreg.predict(X_test)
```

executed in 26ms, finished 00:56:45 2021-04-23

```
In [21]: display(model_log.score(X_train, y_train))
display(model_log.score(X_test, y_test))
```

executed in 14ms, finished 00:56:45 2021-04-23


0.8690922730682671

0.8605697151424287

```
In [22]: # Plot confusion matrix

plot_confusion_matrix(model_log, X_test, y_test, cmap=plt.cm.Blues)
plt.title('Simple Logistic Regression')
plt.show()
```

executed in 171ms, finished 00:56:45 2021-04-23



In [23]: *# we compute our validation metric, recall*

```
print('Training Precision: ', precision_score(y_train, y_hat_train))
print('Testing Precision: ', precision_score(y_test, y_hat_test))
print('\n\n')

print('Training Recall: ', recall_score(y_train, y_hat_train))
print('Testing Recall: ', recall_score(y_test, y_hat_test))
print('\n\n')

print('Training Accuracy: ', accuracy_score(y_train, y_hat_train))
print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
print('\n\n')

print('Training F1-Score: ', f1_score(y_train, y_hat_train))
print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

executed in 27ms, finished 00:56:45 2021-04-23

Training Precision: 0.6211180124223602  
Testing Precision: 0.53125

Training Recall: 0.25773195876288657  
Testing Recall: 0.17894736842105263

Training Accuracy: 0.8690922730682671  
Testing Accuracy: 0.8605697151424287

Training F1-Score: 0.36429872495446264  
Testing F1-Score: 0.2677165354330709

- Accuracy is 86%. The accuracy can possibly be higher on for a different model.
- Testing and testing score can be closer.

## ▼ 1.5.2 Random Forest

In [24]: *# New dataframe for random forest model*  
df\_rf = pd.get\_dummies(df)

executed in 14ms, finished 00:56:45 2021-04-23

In [25]: *# Create X,y and train/test*  
X = df\_dummy.drop(columns=['churn'], axis=1)  
y = df\_dummy['churn']  
  
X\_train, X\_test, y\_train, y\_test = train\_test\_split(X, y, test\_size = 0.25, random\_state=42)

executed in 13ms, finished 00:56:45 2021-04-23



```
In [26]: # Set up initial forest
tree_clf = DecisionTreeClassifier()
tree_clf.fit(X_train, y_train)
```


executed in 45ms, finished 00:56:45 2021-04-23

Out[26]: DecisionTreeClassifier()

```
In [27]: # Plotting feature importance of models
def plot_feature_importances(model):
    n_features = X_train.shape[1]
    plt.figure(figsize=(10,18))
    plt.barh(range(n_features), model.feature_importances_)
    plt.xticks(fontsize=10, fontweight='bold')
    plt.yticks(np.arange(n_features), X_train.columns.values, fontsize=10, fontwe
    plt.xlabel('Importance of Cause', fontsize=16, fontweight='bold' )
    plt.ylabel('Causes', fontsize=16, fontweight='bold')
    plt.grid(linestyle='--', linewidth=2, axis='x', alpha=0.3)
    plt.title("Impact of Factors on Telecom Customer Churning", fontsize=15, font

plot_feature_importances(tree_clf)
```

executed in 1.45s, finished 00:56:47 2021-04-23



```
In [28]: # Test set predictions
pred = tree_clf.predict(X_test)

# Confusion matrix and classification report
print(confusion_matrix(y_test, pred))
print(classification_report(y_test, pred))
print("Testing Accuracy for Decision Tree Classifier: {:.4}%".format(accuracy_score(y_test, pred)))
```

executed in 14ms, finished 00:56:47 2021-04-23


```
[[683  32]
 [ 31  88]]
```

	precision	recall	f1-score	support
0	0.96	0.96	0.96	715
1	0.73	0.74	0.74	119
accuracy			0.92	834
macro avg	0.84	0.85	0.85	834
weighted avg	0.92	0.92	0.92	834

Testing Accuracy for Decision Tree Classifier: 92.45%

```
In [29]: forest = RandomForestClassifier()  
forest.fit(X_train, np.ravel(y_train))  
plot_feature_importances(forest)
```

executed in 1.84s, finished 00:56:48 2021-04-23



- The individual states do not seem to have a large influence on the customer churn
- **Customer\_service\_calls\_** and **total\_day\_charge** appears to have the largest influence on churn

```
In [30]: param_grid = {  
    'max_depth': [2, 5, 10, 25],  
    'min_samples_split': [2, 5, 10, 20]  
}  
  
gs_tree = GridSearchCV(forest, param_grid, cv=3)  
gs_tree.fit(X_train, np.ravel(y_train))  
  
gs_tree.best_params_
```

executed in 10.6s, finished 00:56:59 2021-04-23

```
Out[30]: {'max_depth': 25, 'min_samples_split': 5}
```

```
In [31]: forest = RandomForestClassifier(max_depth=5, min_samples_split=2)
forest.fit(X_train, np.ravel(y_train))
preds = forest.predict(X_test)
print(confusion_matrix(y_test, preds))
print(classification_report(y_test, preds))
print("Testing Accuracy for Random Forest Classifier: {:.4}%".format(accuracy_score(y_test, preds)))
```

executed in 233ms, finished 00:56:59 2021-04-23

```
[[ 715   0]
 [ 119   0]]
```

	precision	recall	f1-score	support
0	0.86	1.00	0.92	715
1	0.00	0.00	0.00	119
accuracy			0.86	834
macro avg	0.43	0.50	0.46	834
weighted avg	0.73	0.86	0.79	834

Testing Accuracy for Random Forest Classifier: 85.73%

C:\Users\leebr\anaconda3\envs\learn-env\lib\site-packages\sklearn\metrics\\_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

- Accuracy is still not very high (85%)
- Should try another model

### ▼ 1.5.3 XG Boost

```
In [32]: # New dataframe for XG Boost model
df_xgb = pd.get_dummies(df)
```

executed in 15ms, finished 00:56:59 2021-04-23

```
In [33]: # Create X,y and train/test
X = df_xgb.drop(columns=['churn'], axis=1)
y = df_xgb['churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state=42)
```

executed in 15ms, finished 00:56:59 2021-04-23

```
In [34]: # Instantiate XGBClassifier
xgb = XGBClassifier()

# Fit XGBClassifier
xgb.fit(X_train, np.ravel(y_train))

# Predict on training and test sets
training_preds = xgb.predict(X_train)
test_preds = xgb.predict(X_test)

# Accuracy of training and test sets
training_accuracy = accuracy_score(y_train, training_preds)
test_accuracy = accuracy_score(y_test, test_preds)

print('Training Accuracy: {:.4}%'.format(training_accuracy * 100))
print('Validation accuracy: {:.4}%'.format(test_accuracy * 100))
```


executed in 514ms, finished 00:57:00 2021-04-23

Training Accuracy: 100.0%  
Validation accuracy: 96.04%



```
In [35]: plot_feature_importances(xgb)
```

executed in 1.57s, finished 00:57:01 2021-04-23





```

In [36]: param_grid = {
    'learning_rate': [0.1, 0.2],
    'max_depth': [6],
    'min_child_weight': [1, 2],
    'subsample': [0.5, 0.7],
    'n_estimators': [100],
}

grid_clf = GridSearchCV(xgb, param_grid, scoring='accuracy')
grid_clf.fit(X_train, np.ravel(y_train))

best_parameters = grid_clf.best_params_

print('Grid Search found the following optimal parameters: ')
for param_name in sorted(best_parameters.keys()):
    print('%s: %r' % (param_name, best_parameters[param_name]))

training_preds = grid_clf.predict(X_train)
test_preds = grid_clf.predict(X_test)
training_accuracy = accuracy_score(y_train, training_preds)
test_accuracy = accuracy_score(y_test, test_preds)

print('')
print('Training Accuracy: {:.4}%'.format(training_accuracy * 100))
print('Validation accuracy: {:.4}%'.format(test_accuracy * 100))

```

executed in 7.09s, finished 00:57:08 2021-04-23

Grid Search found the following optimal parameters:

learning\_rate: 0.1  
 max\_depth: 6  
 min\_child\_weight: 1  
 n\_estimators: 100  
 subsample: 0.7

Training Accuracy: 98.44%  
 Validation accuracy: 96.76%

```

In [37]: print(confusion_matrix(y_test, test_preds))
print(classification_report(y_test, test_preds))
print("Testing Accuracy for XGBoost: {:.4}%".format(accuracy_score(y_test, test_p

```

executed in 30ms, finished 00:57:08 2021-04-23

```

[[712   3]
 [ 24  95]]

```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	715
1	0.97	0.80	0.88	119
accuracy			0.97	834
macro avg	0.97	0.90	0.93	834
weighted avg	0.97	0.97	0.97	834

Testing Accuracy for XGBoost: 96.76%

- XG Boost has the highest testing accuracy from the other models.
- We shall use this for our predictive analysis

### ▼ 1.5.4 Final Model Selection

In [38]: `df_final = pd.DataFrame(xgb.feature_importances_, X.columns, columns = ['coefficient'])`  
`df_final.head()`

executed in 76ms, finished 00:57:09 2021-04-23

Out[38]:

	coefficient
<b>account_length</b>	0.019876
<b>international_plan</b>	0.103945
<b>voice_mail_plan</b>	0.118447
<b>total_day_calls</b>	0.016703
<b>total_day_charge</b>	0.066109

In [39]: `df_final_detail = df_final[:12]`  
`display(df_final_detail.head())`  
`df_final_state = df_final[12:]`  
`display(df_final_state.head())`

executed in 14ms, finished 00:57:09 2021-04-23

	coefficient
<b>account_length</b>	0.019876
<b>international_plan</b>	0.103945
<b>voice_mail_plan</b>	0.118447
<b>total_day_calls</b>	0.016703
<b>total_day_charge</b>	0.066109

	coefficient
<b>state_AK</b>	0.000000
<b>state_AL</b>	0.012708
<b>state_AR</b>	0.000000
<b>state_AZ</b>	0.007599
<b>state_CA</b>	0.011509

```
In [40]: # Renaming states indexes
as_list = df_final_state.index.values.tolist()
as_list = [i.replace('state_', '') for i in as_list]
df_final_state.index = as_list
df_final_state.head()
```

executed in 15ms, finished 00:57:09 2021-04-23


Out[40]:

	<b>coefficient</b>
<b>AK</b>	0.000000
<b>AL</b>	0.012708
<b>AR</b>	0.000000
<b>AZ</b>	0.007599
<b>CA</b>	0.011509

```
In [41]: # Create a barplot of state's impact on churn
ax = df_final_state.sort_values(by=['coefficient'], ascending=False).plot(kind='bar')
ax.set_title("Impact of States on Customer Churning", fontsize = 22, fontweight = 'bold')
ax.set_xlabel("State", fontsize=20, fontweight='bold')
ax.set_ylabel("Impact on Customer Churn decision", fontsize=16, fontweight='bold')
ax.get_legend().remove()

plt.xticks(rotation=-65, fontsize=12, fontweight='bold')
plt.yticks(fontsize=12, fontweight='bold')
plt.grid(color='#95a5a6', linestyle='--', linewidth=2, axis='y', alpha=0.3)
```

executed in 1.06s, finished 00:57:10 2021-04-23



- There seems to be a large churn rate in **state\_OR**. This may need to be investigated further.
- No apparent correlation can be found with states with high impact

```
In [42]: # Renaming detail indexes
detail_list = df_final_detail.index.values.tolist()
detail_list = [i.replace('_', ' ') for i in detail_list]
df_final_detail.index = detail_list
df_final_detail.head()
```

executed in 14ms, finished 00:57:10 2021-04-23

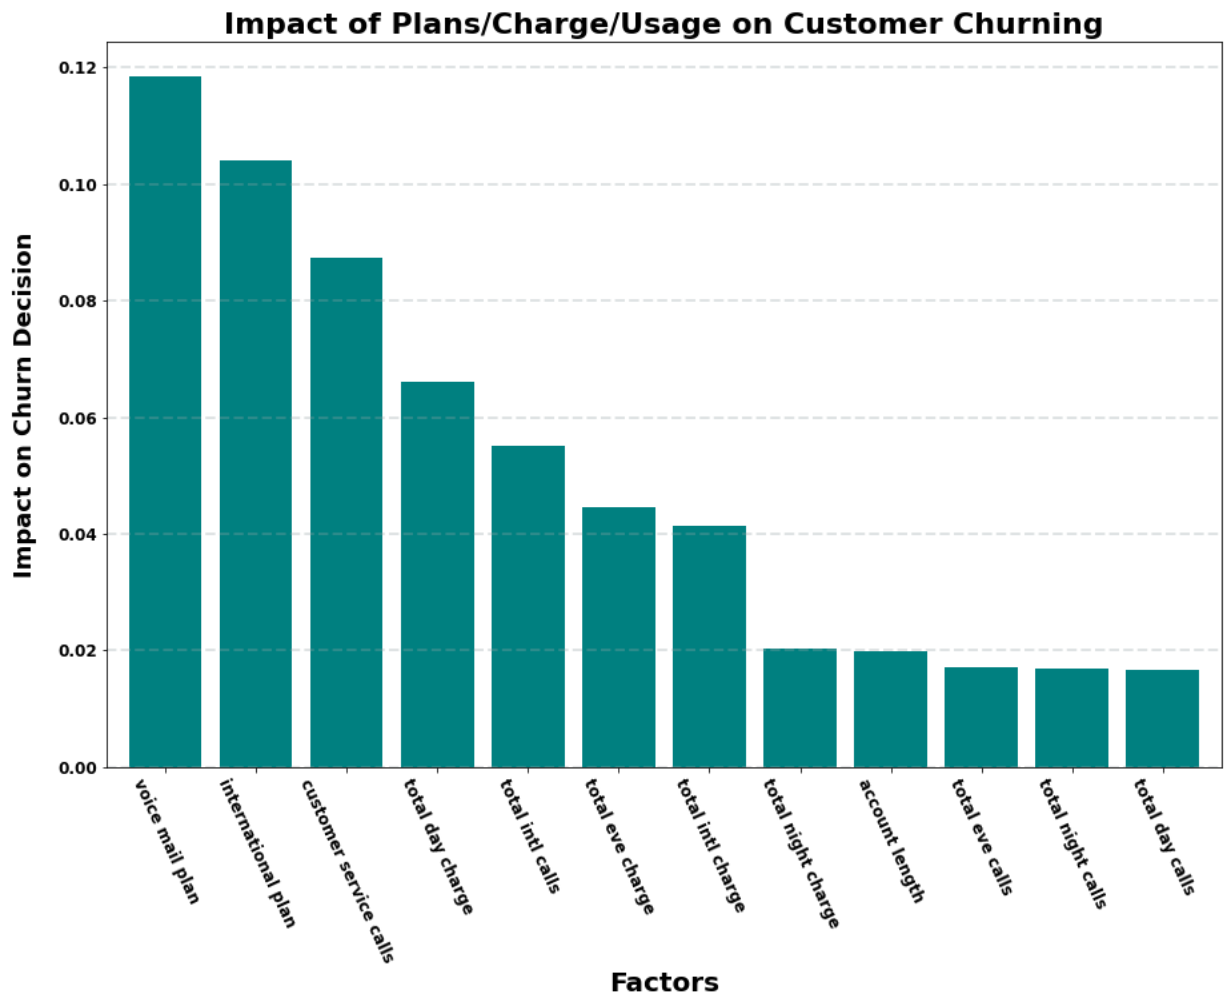
Out[42]:

	coefficient
account length	0.019876
international plan	0.103945
voice mail plan	0.118447
total day calls	0.016703
total day charge	0.066109

```
In [43]: # Create a barplot of customer plans/usage's impact on churn
ax = df_final_detail.sort_values(by=['coefficient'], ascending=False).plot(kind='bar')
ax.set_title("Impact of Plans/Charge/Usage on Customer Churning", fontsize = 22,
ax.set_xlabel("Factors", fontsize=20, fontweight='bold')
ax.set_ylabel("Impact on Churn Decision", fontsize=18, fontweight='bold', labelpad=10)
ax.get_legend().remove()

plt.xticks(rotation=-65, fontsize=12, fontweight='bold')
plt.yticks(fontsize=12, fontweight='bold')
plt.grid(color='#95a5a6', linestyle='--', linewidth=2, axis='y', alpha=0.3)
```

executed in 251ms, finished 00:57:10 2021-04-23



- **voice\_mail\_plan** and **international\_plan** are the largest factors to customer churn. There may be some issues with how both plans are being handled.
- **customer\_service\_calls** come close in influence to the plans. There may need to be an improvement to the customer service section of the company.

## 1.6 Conclusions

The analysis of the SyriaTel customer churn dataset resulted in the following conclusions:



- There needs to be an improvement to the **Voice mail plan** and **International plan**. The customer churn is heavily affected by the effectiveness of the plans. These plans need to be further examined to entice customers to stay with SyriaTel.
- The **Customer service** department may need examining. We need to check staffing to see what is causing the customers to turn away from SyriaTel as they submit for help through the customer service line.
- We can predict future customer churn with our final model. This should help us mitigate customer losses if we contact the customer earlier for their input.

## ▼ 1.7 Next Steps

Further analysis of the SyriaTel data could yield additional insights to other recommendations

- **Locate what factors are causing a larger churn within specific states**
- **Create an alert system that detects when individual customer are in range of possibly churning**
- **Investigate a change over system from international to domestic plans on customer churn**